

avacado

June 18, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df=pd.read_csv("C:/Users/Ravi/Downloads/MLR/Datasets_MLR/Avacado_Price.csv")
```

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   AveragePrice    18249 non-null  float64
 1   Total_Volume    18249 non-null  float64
 2   tot_ava1        18249 non-null  float64
 3   tot_ava2        18249 non-null  float64
 4   tot_ava3        18249 non-null  float64
 5   Total_Bags      18249 non-null  float64
 6   Small_Bags      18249 non-null  float64
 7   Large_Bags      18249 non-null  float64
 8   Xlarge_Bags     18249 non-null  float64
 9   type            18249 non-null  object
10   year            18249 non-null  int64
11   region          18249 non-null  object
dtypes: float64(9), int64(1), object(2)
memory usage: 1.7+ MB
```

```
[4]: duplicated=df.duplicated()
```

```
[5]: sum(duplicated)
```

```
[5]: 0
```

```
[6]: df.isna().sum()
```

```
[6]: AveragePrice    0
      Total_Volume   0
      tot_ava1       0
      tot_ava2       0
      tot_ava3       0
      Total_Bags     0
      Small_Bags     0
      Large_Bags     0
      Xlarge_Bags    0
      type           0
      year           0
      region         0
      dtype: int64
```

```
[7]: df.isna().sum()
```

```
[7]: AveragePrice    0
      Total_Volume   0
      tot_ava1       0
      tot_ava2       0
      tot_ava3       0
      Total_Bags     0
      Small_Bags     0
      Large_Bags     0
      Xlarge_Bags    0
      type           0
      year           0
      region         0
      dtype: int64
```

```
[8]: df.shape
```

```
[8]: (18249, 12)
```

```
[9]: df.head()
```

```
[9]:   AveragePrice  Total_Volume  tot_ava1  tot_ava2  tot_ava3  Total_Bags  \
0         1.33      64236.62   1036.74   54454.85      48.16      8696.87
1         1.35      54876.98    674.28   44638.81      58.33      9505.56
2         0.93     118220.22    794.70  109149.67     130.50      8145.35
3         1.08      78992.15   1132.00   71976.41      72.58      5811.16
4         1.28      51039.60    941.48   43838.39      75.78      6183.95

      Small_Bags  Large_Bags  Xlarge_Bags      type  year  region
0      8603.62     93.25         0.0  conventional  2015  Albany
1      9408.07     97.49         0.0  conventional  2015  Albany
2      8042.21    103.14         0.0  conventional  2015  Albany
3      5677.40    133.76         0.0  conventional  2015  Albany
```

4 5986.26 197.69 0.0 conventional 2015 Albany

```
[10]: from sklearn.preprocessing import LabelEncoder
      le=LabelEncoder()
```

```
[11]: df.type=le.fit_transform(df.type)
```

```
[12]: df.region=le.fit_transform(df.region)
```

```
[13]: df.describe()
```

```
[13]:
```

	AveragePrice	Total_Volume	tot_ava1	tot_ava2	tot_ava3	\
count	18249.000000	1.824900e+04	1.824900e+04	1.824900e+04	1.824900e+04	
mean	1.405978	8.506440e+05	2.930084e+05	2.951546e+05	2.283974e+04	
std	0.402677	3.453545e+06	1.264989e+06	1.204120e+06	1.074641e+05	
min	0.440000	8.456000e+01	0.000000e+00	0.000000e+00	0.000000e+00	
25%	1.100000	1.083858e+04	8.540700e+02	3.008780e+03	0.000000e+00	
50%	1.370000	1.073768e+05	8.645300e+03	2.906102e+04	1.849900e+02	
75%	1.660000	4.329623e+05	1.110202e+05	1.502069e+05	6.243420e+03	
max	3.250000	6.250565e+07	2.274362e+07	2.047057e+07	2.546439e+06	

	Total_Bags	Small_Bags	Large_Bags	Xlarge_Bags	type	\
count	1.824900e+04	1.824900e+04	1.824900e+04	18249.000000	18249.000000	
mean	2.396392e+05	1.821947e+05	5.433809e+04	3106.426507	0.499918	
std	9.862424e+05	7.461785e+05	2.439660e+05	17692.894652	0.500014	
min	0.000000e+00	0.000000e+00	0.000000e+00	0.000000	0.000000	
25%	5.088640e+03	2.849420e+03	1.274700e+02	0.000000	0.000000	
50%	3.974383e+04	2.636282e+04	2.647710e+03	0.000000	0.000000	
75%	1.107834e+05	8.333767e+04	2.202925e+04	132.500000	1.000000	
max	1.937313e+07	1.338459e+07	5.719097e+06	551693.650000	1.000000	

	year	region
count	18249.000000	18249.000000
mean	2016.147899	26.495644
std	0.939938	15.583788
min	2015.000000	0.000000
25%	2015.000000	13.000000
50%	2016.000000	26.000000
75%	2017.000000	40.000000
max	2018.000000	53.000000

```
[75]: plt.bar(height=df.AveragePrice,x=np.arange(1,18250,1));plt.
      ↪title("      AveragePrice")
```


```
[75]: Text(0.5, 1.0, '\tAveragePrice')
```

C:\Users\Ravi\anaconda3\lib\site-packages\matplotlib\backends\backend_agg.py:240: RuntimeWarning: Glyph 9 missing

```

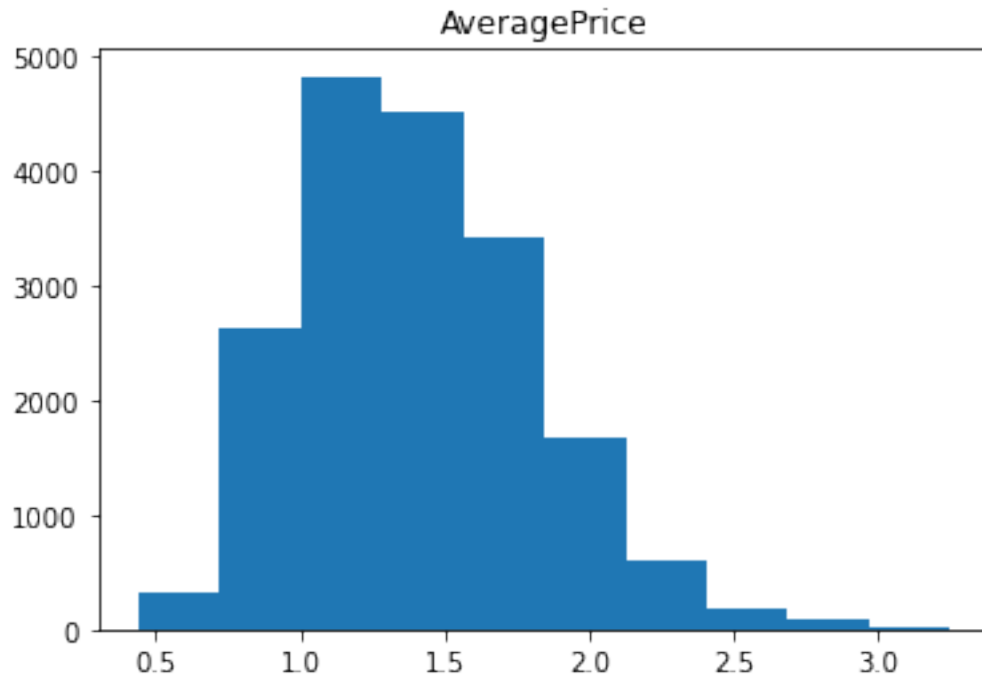
from current font.
    font.set_text(s, 0.0, flags=flags)
C:\Users\Ravi\anaconda3\lib\site-
packages\matplotlib\backends\backend_agg.py:203: RuntimeWarning: Glyph 9 missing
from current font.
    font.set_text(s, 0, flags=flags)

```




```
[76]: plt.hist(df.AveragePrice);plt.title("AveragePrice")
```

```
[76]: Text(0.5, 1.0, 'AveragePrice')
```




```
[145]: plt.boxplot(df.AveragePrice)
```

```
[145]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e50d424f0>,
<matplotlib.lines.Line2D at 0x12e50cb8820>],
'caps': [<matplotlib.lines.Line2D at 0x12e50cc4760>,
<matplotlib.lines.Line2D at 0x12e509faa60>],
'boxes': [<matplotlib.lines.Line2D at 0x12e4641f610>],
'medians': [<matplotlib.lines.Line2D at 0x12e50d658e0>],
'fliers': [<matplotlib.lines.Line2D at 0x12e50d65bb0>],
'means': []}
```



```
[77]: plt.bar(height=df.Total_Volume,x=np.arange(1,18250,1));plt.title("Total_Volume")
```

```
[77]: Text(0.5, 1.0, 'Total_Volume')
```




```
[146]: IQR=df.AveragePrice.quantile(0.75)-df.AveragePrice.quantile(0.25)
lower_limit=df.Total_Volume.quantile(0.25)-(IQR*1.5)
upper_limit=df.Total_Volume.quantile(0.75)-(IQR*1.5)

outliers=np.where(df.AveragePrice>upper_limit,True,np.where(df.
    ↪AveragePrice<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↪gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['AveragePrice'])


[147]: df_t=winsor.fit_transform(df[["AveragePrice"]])
plt.boxplot(df_t)
```

```
[147]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e50d8f820>,
    <matplotlib.lines.Line2D at 0x12e50d8fbb0>],
'caps': [<matplotlib.lines.Line2D at 0x12e50d8ff40>,
    <matplotlib.lines.Line2D at 0x12e50dbb310>],
'boxes': [<matplotlib.lines.Line2D at 0x12e50d8f490>],
'medians': [<matplotlib.lines.Line2D at 0x12e50dbb6a0>],
'fliers': [<matplotlib.lines.Line2D at 0x12e50dbba30>],
'means': []}
```




```
[78]: plt.hist(df.Total_Volume);plt.title("Total_Volume")
```

```
[78]: Text(0.5, 1.0, 'Total_Volume')
```



```
[79]: plt.boxplot(df.Total_Volume)
```

```
[79]: {'whiskers': [<matplotlib.lines.Line2D at 0x12ddc0f3130>,
<matplotlib.lines.Line2D at 0x12ddc0f34c0>],
'caps': [<matplotlib.lines.Line2D at 0x12ddc0f3850>,
<matplotlib.lines.Line2D at 0x12ddc0f3be0>],
'boxes': [<matplotlib.lines.Line2D at 0x12ddc0e6d60>],
'medians': [<matplotlib.lines.Line2D at 0x12ddc0f3f70>],
'fliers': [<matplotlib.lines.Line2D at 0x12ddc0fd340>],
'means': []}
```


```
[80]: IQR=df.Total_Volume.quantile(0.75)-df.Total_Volume.quantile(0.25)
lower_limit=df.Total_Volume.quantile(0.25)-(IQR*1.5)
upper_limit=df.Total_Volume.quantile(0.75)-(IQR*1.5)
```

```
[81]: outliers=np.where(df.Total_Volume>upper_limit,True,np.where(df.
    ↪Total_Volume<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↪gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['Total_Volume'])
```

```
[82]: df_t=winsor.fit_transform(df[["Total_Volume"]])
```

```
[83]: plt.boxplot(df_t)
```


```
[83]: {'whiskers': [<matplotlib.lines.Line2D at 0x12ddc15c190>,
    <matplotlib.lines.Line2D at 0x12ddc15c520>],
'caps': [<matplotlib.lines.Line2D at 0x12ddc15c8b0>,
    <matplotlib.lines.Line2D at 0x12ddc15cc40>],
'boxes': [<matplotlib.lines.Line2D at 0x12ddc14cdc0>],
'medians': [<matplotlib.lines.Line2D at 0x12ddc15cfd0>],
'fliers': [<matplotlib.lines.Line2D at 0x12ddc1653a0>],
'means': []}
```



```
[84]: df.Total_Volume=df_t
```


```
[85]: plt.boxplot(df.Total_Volume)
```

```
[85]: {'whiskers': [<matplotlib.lines.Line2D at 0x12ddc1c6100>,
<matplotlib.lines.Line2D at 0x12ddc1c6490>],
'caps': [<matplotlib.lines.Line2D at 0x12ddc1c6820>,
<matplotlib.lines.Line2D at 0x12ddc1c6bb0>],
'boxes': [<matplotlib.lines.Line2D at 0x12ddc1b7d30>],
'medians': [<matplotlib.lines.Line2D at 0x12ddc1c6f40>],
'fliers': [<matplotlib.lines.Line2D at 0x12ddc1d0310>],
'means': []}
```




```
[86]: plt.hist(df.Total_Volume);plt.title("Total_Volume")
```

```
[86]: Text(0.5, 1.0, 'Total_Volume')
```




```
[87]: plt.bar(height=df.Total_Volume,x=np.arange(1,18250,1));plt.title("Total_Volume")
```

```
[87]: Text(0.5, 1.0, 'Total_Volume')
```




```
[88]: sns.jointplot(x=df.Total_Volume,y=df.AveragePrice)
```

```
[88]: <seaborn.axisgrid.JointGrid at 0x12ddc24a160>
```




```
[89]: plt.bar(height=df.tot_ava1,x=np.arange(1,18250,1));plt.title("Tot_ava1")
```

```
[89]: Text(0.5, 1.0, 'Tot_ava1')
```



```
[90]: plt.boxplot(df.tot_ava1)
```

```
[90]: {'whiskers': [<matplotlib.lines.Line2D at 0x12df70b8a30>,
<matplotlib.lines.Line2D at 0x12df70b8dc0>],
'caps': [<matplotlib.lines.Line2D at 0x12df70c7190>,
<matplotlib.lines.Line2D at 0x12df70c7520>],
'boxes': [<matplotlib.lines.Line2D at 0x12df70b86a0>],
'medians': [<matplotlib.lines.Line2D at 0x12df70c78b0>],
'fliers': [<matplotlib.lines.Line2D at 0x12df70c7c40>],
'means': []}
```



```
[91]: IQR=df.tot_ava1.quantile(0.75)-df.tot_ava1.quantile(0.25)
lower_limit=df.tot_ava1.quantile(0.25)-(IQR*1.5)
upper_limit=df.tot_ava1.quantile(0.75)-(IQR*1.5)

outliers=np.where(df.tot_ava1>upper_limit,True,np.where(df.
    ↪tot_ava1<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↪gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['tot_ava1'])
```


```
[92]: df_t=winsor.fit_transform(df[["tot_ava1"]])
```

```
[93]: df.tot_ava1=df_t
```

```
[94]: plt.boxplot(df.tot_ava1)
```

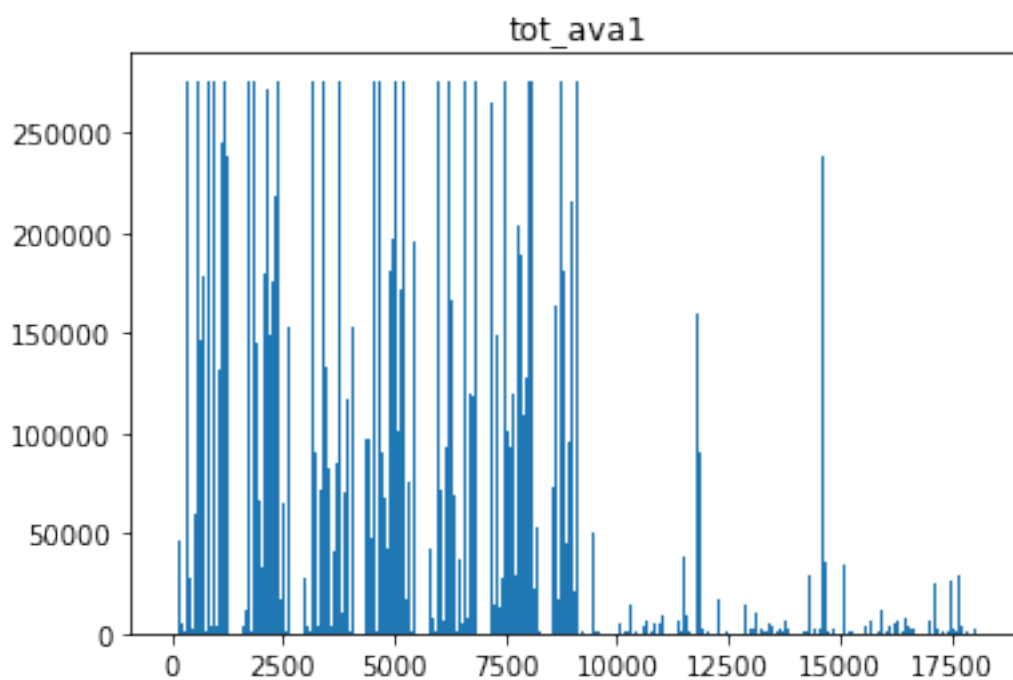
```
[94]: {'whiskers': [<matplotlib.lines.Line2D at 0x12df6723760>,
    <matplotlib.lines.Line2D at 0x12df6723af0>],
    'caps': [<matplotlib.lines.Line2D at 0x12df6723e80>,
    <matplotlib.lines.Line2D at 0x12df672d250>],
    'boxes': [<matplotlib.lines.Line2D at 0x12df67233d0>],
```

```
'medians': [<matplotlib.lines.Line2D at 0x12df672d5e0>],  
'fliers': [<matplotlib.lines.Line2D at 0x12df672d970>],  
'means': []}
```



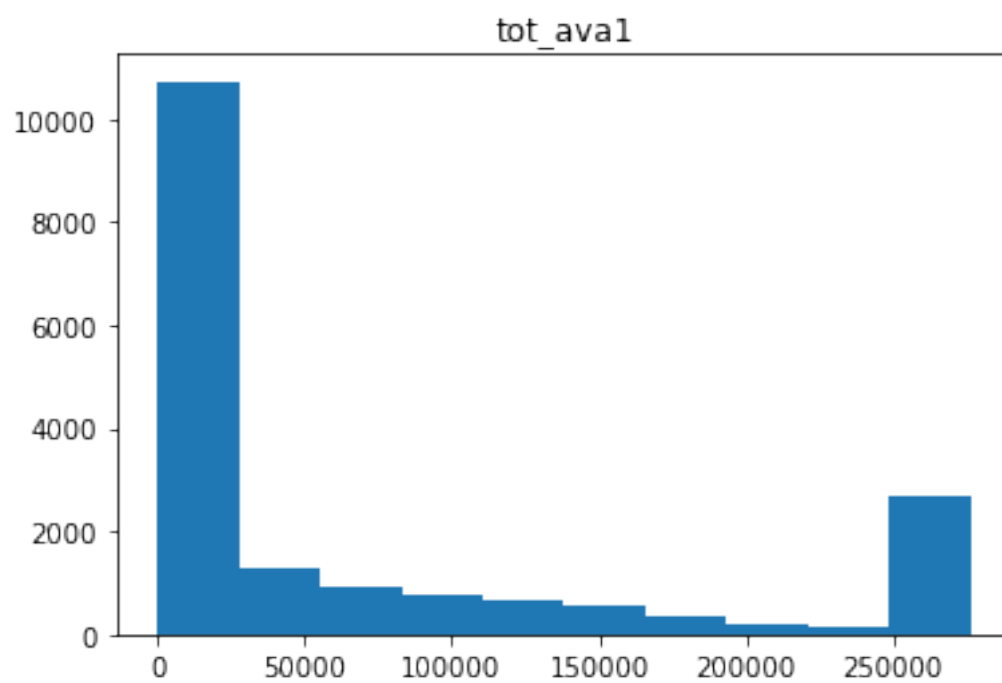
```
[95]: plt.bar(height=df.tot_ava1,x=np.arange(1,18250,1));plt.title("tot_ava1")
```

```
[95]: Text(0.5, 1.0, 'tot_ava1')
```


```
[96]: plt.hist(df.tot_ava1);plt.title("tot_ava1")
```

```
[96]: Text(0.5, 1.0, 'tot_ava1')
```



```
[97]: sns.jointplot(x=df.tot_ava1,y=df.AveragePrice)
```


```
[97]: <seaborn.axisgrid.JointGrid at 0x12de9557040>
```



```
[98]: plt.boxplot(df.tot_ava2)
```

```
[98]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e03b13a90>,  
                 <matplotlib.lines.Line2D at 0x12e03b13e20>],  
      'caps': [<matplotlib.lines.Line2D at 0x12e03b221f0>,  
              <matplotlib.lines.Line2D at 0x12e03b22580>],  
      'boxes': [<matplotlib.lines.Line2D at 0x12e03b13700>],
```

```
'medians': [<matplotlib.lines.Line2D at 0x12e03b22910>],
'fliers': [<matplotlib.lines.Line2D at 0x12e03b22ca0>],
'means': []}
```



```
[99]: IQR=df.tot_ava2.quantile(0.75)-df.tot_ava2.quantile(0.25)
lower_limit=df.tot_ava2.quantile(0.25)-(IQR*1.5)
upper_limit=df.tot_ava2.quantile(0.75)-(IQR*1.5)


outliers=np.where(df.tot_ava2>upper_limit,True,np.where(df.
    ↳tot_ava2<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↳gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['tot_ava2'])
```

```
[100]: df_t=winsor.fit_transform(df[["tot_ava2"]])
```

```
[101]: df.tot_ava2=df_t
```


```
[102]: plt.boxplot(df.tot_ava2);plt.title("tot_ava2")
```

```
[102]: Text(0.5, 1.0, 'tot_ava2')
```




```
[103]: plt.bar(height=df.tot_ava2,x=np.arange(1,18250,1));plt.title("tot_ava2")
```

```
[103]: Text(0.5, 1.0, 'tot_ava2')
```




```
[104]: plt.hist(df.tot_ava2);plt.title("tot_ava2")
```

```
[104]: Text(0.5, 1.0, 'tot_ava2')
```




```
[105]: sns.jointplot(x=df.tot_ava2,y=df.AveragePrice)
```

```
[105]: <seaborn.axisgrid.JointGrid at 0x12e03b37850>
```



```
[106]: plt.boxplot(df.tot_ava3)
```

```
[106]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e10f91e80>,
<matplotlib.lines.Line2D at 0x12e10fa0250>],
'caps': [<matplotlib.lines.Line2D at 0x12e10fa05e0>,
<matplotlib.lines.Line2D at 0x12e10fa0970>],
'boxes': [<matplotlib.lines.Line2D at 0x12e10f91af0>],
'medians': [<matplotlib.lines.Line2D at 0x12e10fa0d00>],
'fliers': [<matplotlib.lines.Line2D at 0x12e10fad0d0>],
'means': []}
```



```
[107]: IQR=df.tot_ava3.quantile(0.75)-df.tot_ava3.quantile(0.25)
lower_limit=df.tot_ava3.quantile(0.25)-(IQR*1.5)
upper_limit=df.tot_ava3.quantile(0.75)-(IQR*1.5)

outliers=np.where(df.tot_ava3>upper_limit,True,np.where(df.
    ↪tot_ava3<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↪gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['tot_ava3'])
```


```
[108]: df_t=winsor.fit_transform(df[["tot_ava3"]])
```

```
[109]: df.tot_ava3=df_t
```

```
[110]: plt.boxplot(df.tot_ava3)
```


```
[110]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e11003c10>,
    <matplotlib.lines.Line2D at 0x12e11003fa0>],
    'caps': [<matplotlib.lines.Line2D at 0x12e11010370>,
    <matplotlib.lines.Line2D at 0x12e11010700>],
    'boxes': [<matplotlib.lines.Line2D at 0x12e11003880>],
```

```
'medians': [<matplotlib.lines.Line2D at 0x12e11010a90>],  
'fliers': [<matplotlib.lines.Line2D at 0x12e11010e20>],  
'means': []}
```




```
[111]: plt.bar(height=df.tot_ava3,x=np.arange(1,18250,1));plt.title("tot_ava3")
```

```
[111]: Text(0.5, 1.0, 'tot_ava3')
```


```
[112]: plt.hist(df.tot_ava3);plt.title("tot_ava3")
```

```
[112]: Text(0.5, 1.0, 'tot_ava3')
```



```
[113]: plt.boxplot(df.Total_Bags)
```

```
[113]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e1ec4abe0>,  
                  <matplotlib.lines.Line2D at 0x12e1ec4af70>],  
        'caps': [<matplotlib.lines.Line2D at 0x12e1e259340>,  
                 <matplotlib.lines.Line2D at 0x12e1e2596d0>],  
        'boxes': [<matplotlib.lines.Line2D at 0x12e1ec4a820>],  
        'medians': [<matplotlib.lines.Line2D at 0x12e1e259a60>],  
        'fliers': [<matplotlib.lines.Line2D at 0x12e1e259e20>],  
        'means': []}
```




```
[114]: IQR=df.Total_Bags.quantile(0.75)-df.Total_Bags.quantile(0.25)  
lower_limit=df.Total_Bags.quantile(0.25)-(IQR*1.5)  
upper_limit=df.Total_Bags.quantile(0.75)-(IQR*1.5)  
  
outliers=np.where(df.Total_Bags>upper_limit,True,np.where(df.  
    ↳Total_Bags<lower_limit,True,False))  
from feature_engine.outliers import Winsorizer  
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or  
    ↳gaussian for mean and std  
                    tail='both', # cap left, right or both tails  
                    fold=1.5,  
                    variables=['Total_Bags'])
```

```
[115]: df_t=winsor.fit_transform(df[["Total_Bags"]])
```

```
[116]: df.Total_Bags=df_t
```

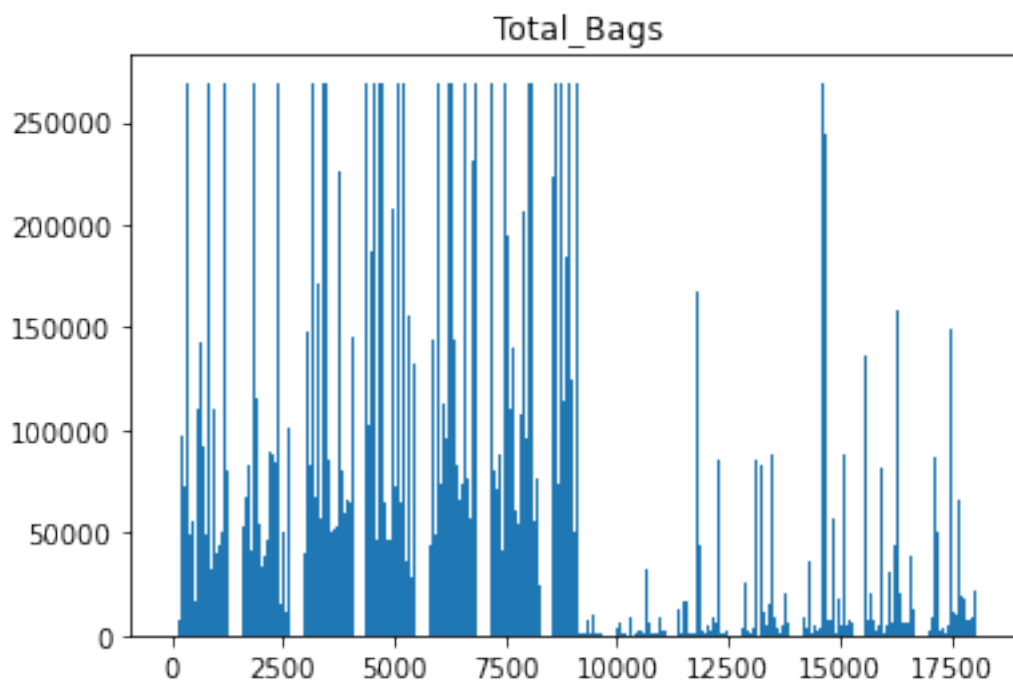
```
[117]: plt.boxplot(df.Total_Bags)
```

```
[117]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e1e2b4910>,  
                  <matplotlib.lines.Line2D at 0x12e1e2b4ca0>],  
       'caps': [<matplotlib.lines.Line2D at 0x12e1e2c2070>,  
               <matplotlib.lines.Line2D at 0x12e1e2c2400>],  
       'boxes': [<matplotlib.lines.Line2D at 0x12e1e2b4580>],  
       'medians': [<matplotlib.lines.Line2D at 0x12e1e2c2790>],  
       'fliers': [<matplotlib.lines.Line2D at 0x12e1e2c2b20>],  
       'means': []}
```



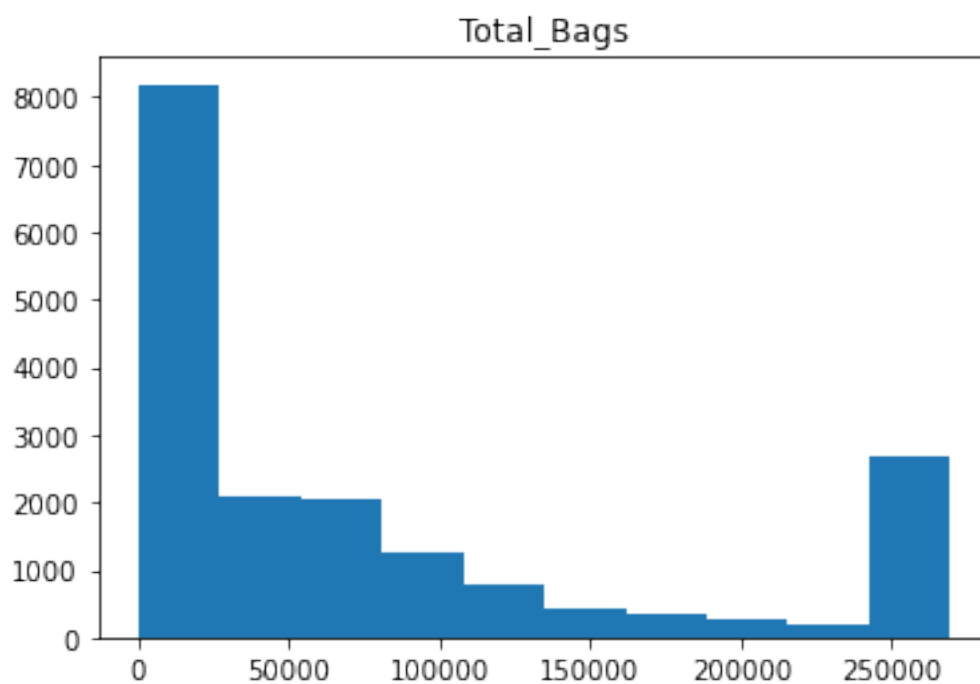
```
[118]: plt.bar(height=df.Total_Bags,x=np.arange(1,18250,1));plt.title("Total_Bags")
```

```
[118]: Text(0.5, 1.0, 'Total_Bags')
```




```
[119]: plt.hist(df.Total_Bags);plt.title("Total_Bags")
```

```
[119]: Text(0.5, 1.0, 'Total_Bags')
```



```
[120]: sns.jointplot(x=df.Total_Bags,y=df.AveragePrice)
```


```
[120]: <seaborn.axisgrid.JointGrid at 0x12e2b4b32b0>
```



```
[121]: plt.boxplot(df.Small_Bags)
```

```
[121]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e2b698940>,  
<matplotlib.lines.Line2D at 0x12e2b698cd0>],  
'caps': [<matplotlib.lines.Line2D at 0x12e2b6a60a0>,  
<matplotlib.lines.Line2D at 0x12e2b6a6430>],  
'boxes': [<matplotlib.lines.Line2D at 0x12e2b6985b0>],
```

```
'medians': [<matplotlib.lines.Line2D at 0x12e2b6a67c0>],
'fliers': [<matplotlib.lines.Line2D at 0x12e2b6a6b50>],
'means': []}
```



```
[122]: IQR=df.Small_Bags.quantile(0.75)-df.Small_Bags.quantile(0.25)
lower_limit=df.Small_Bags.quantile(0.25)-(IQR*1.5)
upper_limit=df.Small_Bags.quantile(0.75)-(IQR*1.5)

outliers=np.where(df.Small_Bags>upper_limit,True,np.where(df.
    ↪Small_Bags<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↪gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['Small_Bags'])
```

```
[123]: df_t=winsor.fit_transform(df[["Small_Bags"]])
```

```
[124]: df.Small_Bags=df_t
```


```
[125]: plt.boxplot(df.Small_Bags)
```

```
[125]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e2b707820>,
<matplotlib.lines.Line2D at 0x12e2b707bb0>],
'caps': [<matplotlib.lines.Line2D at 0x12e2b707f40>,
<matplotlib.lines.Line2D at 0x12e2b712310>],
'boxes': [<matplotlib.lines.Line2D at 0x12e2b707490>],
'medians': [<matplotlib.lines.Line2D at 0x12e2b7126a0>],
'fliers': [<matplotlib.lines.Line2D at 0x12e2b712a30>],
'means': []}
```




```
[126]: plt.bar(height=df.Small_Bags,x=np.arange(1,18250,1));plt.title("Small_Bags")
```

```
[126]: Text(0.5, 1.0, 'Small_Bags')
```




```
[127]: plt.hist(df.Small_Bags);plt.title("Small_Bags")
```

```
[127]: Text(0.5, 1.0, 'Small_Bags')
```




```
[128]: sns.jointplot(x=df.Small_Bags,y=df.AveragePrice)
```


```
[128]: <seaborn.axisgrid.JointGrid at 0x12e392988b0>
```



```
[129]: plt.boxplot(df.Large_Bags)
```

```
[129]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e38afcaf0>,  
                  <matplotlib.lines.Line2D at 0x12e38afce80>],  
       'caps': [<matplotlib.lines.Line2D at 0x12e38b0a250>,  
                <matplotlib.lines.Line2D at 0x12e38b0a5e0>],  
       'boxes': [<matplotlib.lines.Line2D at 0x12e38afc760>],
```

```
'medians': [<matplotlib.lines.Line2D at 0x12e38b0a970>],
'fliers': [<matplotlib.lines.Line2D at 0x12e38b0ad00>],
'means': []}
```



```
[130]: IQR=df.Large_Bags.quantile(0.75)-df.Large_Bags.quantile(0.25)
lower_limit=df.Large_Bags.quantile(0.25)-(IQR*1.5)
upper_limit=df.Large_Bags.quantile(0.75)-(IQR*1.5)


outliers=np.where(df.Large_Bags>upper_limit,True,np.where(df.
    ↪Large_Bags<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↪gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['Large_Bags'])
```

```
[131]: df_t=winsor.fit_transform(df[["Large_Bags"]])
```

```
[132]: df.Large_Bags=df_t
```


```
[133]: plt.boxplot(df.Large_Bags)
```

```
[133]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e38b67760>,
  <matplotlib.lines.Line2D at 0x12e38b67af0>],
  'caps': [<matplotlib.lines.Line2D at 0x12e38b67e80>,
  <matplotlib.lines.Line2D at 0x12e38b72250>],
  'boxes': [<matplotlib.lines.Line2D at 0x12e38b673d0>],
  'medians': [<matplotlib.lines.Line2D at 0x12e38b725e0>],
  'fliers': [<matplotlib.lines.Line2D at 0x12e38b72970>],
  'means': []}
```




```
[134]: plt.bar(height=df.Large_Bags,x=np.arange(1,18250,1));plt.title("Large_Bags")
```

```
[134]: Text(0.5, 1.0, 'Large_Bags')
```




```
[135]: plt.hist(df.Large_Bags);plt.title("Large_Bags")
```

```
[135]: Text(0.5, 1.0, 'Large_Bags')
```



```
[136]: plt.boxplot(df.Xlarge_Bags)
```

```
[136]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e4775bcd0>,
<matplotlib.lines.Line2D at 0x12e45d980a0>],
'caps': [<matplotlib.lines.Line2D at 0x12e45d98430>,
<matplotlib.lines.Line2D at 0x12e45d987c0>],
'boxes': [<matplotlib.lines.Line2D at 0x12e4775b910>],
'medians': [<matplotlib.lines.Line2D at 0x12e45d98b50>],
'fliers': [<matplotlib.lines.Line2D at 0x12e45d98ee0>],
'means': []}
```



```
[137]: IQR=df.Large_Bags.quantile(0.75)-df.Xlarge_Bags.quantile(0.25)
lower_limit=df.Xlarge_Bags.quantile(0.25)-(IQR*1.5)
upper_limit=df.Xlarge_Bags.quantile(0.75)-(IQR*1.5)


outliers=np.where(df.Xlarge_Bags>upper_limit,True,np.where(df.
    ↳Xlarge_Bags<lower_limit,True,False))
from feature_engine.outliers import Winsorizer
winsor = Winsorizer(capping_method='iqr', # choose IQR rule boundaries or
    ↳gaussian for mean and std
                    tail='both', # cap left, right or both tails
                    fold=1.5,
                    variables=['Xlarge_Bags'])
```

```
[138]: df_t=winsor.fit_transform(df[["Xlarge_Bags"]])
```

```
[139]: df.Xlarge_Bags=df_t
```

```
[140]: plt.boxplot(df.Xlarge_Bags)
```

```
[140]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e45df2af0>,  
                  <matplotlib.lines.Line2D at 0x12e45df2e80>],  
       'caps': [<matplotlib.lines.Line2D at 0x12e45dff250>,  
               <matplotlib.lines.Line2D at 0x12e45dff5e0>],  
       'boxes': [<matplotlib.lines.Line2D at 0x12e45df2760>],  
       'medians': [<matplotlib.lines.Line2D at 0x12e45dff970>],  
       'fliers': [<matplotlib.lines.Line2D at 0x12e45dff00>],  
       'means': []}
```



```
[148]: plt.boxplot(df)
```

```
[148]: {'whiskers': [<matplotlib.lines.Line2D at 0x12e50df9340>,  
                  <matplotlib.lines.Line2D at 0x12e50df96d0>,  
                  <matplotlib.lines.Line2D at 0x12e50e02c70>,  
                  <matplotlib.lines.Line2D at 0x12e50e0e040>,  
                  <matplotlib.lines.Line2D at 0x12e51f485e0>,  
                  <matplotlib.lines.Line2D at 0x12e51f48970>,  
                  <matplotlib.lines.Line2D at 0x12e51f56f10>,  
                  <matplotlib.lines.Line2D at 0x12e51f612e0>,  
                  <matplotlib.lines.Line2D at 0x12e51f6c880>],  
       'caps': [...],  
       'boxes': [...],  
       'medians': [...],  
       'fliers': [...],  
       'means': [...]}
```

```


<matplotlib.lines.Line2D at 0x12e51f6cc10>,
<matplotlib.lines.Line2D at 0x12e51f821f0>,
<matplotlib.lines.Line2D at 0x12e51f82580>,
<matplotlib.lines.Line2D at 0x12e51f8eb20>,
<matplotlib.lines.Line2D at 0x12e51f8eeb0>,
<matplotlib.lines.Line2D at 0x12e51fa4490>,
<matplotlib.lines.Line2D at 0x12e51fa4820>,
<matplotlib.lines.Line2D at 0x12e51faedc0>,
<matplotlib.lines.Line2D at 0x12e51fba190>,
<matplotlib.lines.Line2D at 0x12e51fc5730>,
<matplotlib.lines.Line2D at 0x12e51fc5ac0>,
<matplotlib.lines.Line2D at 0x12e51fdc0a0>,
<matplotlib.lines.Line2D at 0x12e51fdc430>,
<matplotlib.lines.Line2D at 0x12e51fe69d0>,
<matplotlib.lines.Line2D at 0x12e51fe6d60>],
'caps': [<matplotlib.lines.Line2D at 0x12e50df9a60>,
<matplotlib.lines.Line2D at 0x12e50df9df0>,
<matplotlib.lines.Line2D at 0x12e50e0e3d0>,
<matplotlib.lines.Line2D at 0x12e50e0e760>,
<matplotlib.lines.Line2D at 0x12e51f48d00>,
<matplotlib.lines.Line2D at 0x12e51f560d0>,
<matplotlib.lines.Line2D at 0x12e51f61670>,
<matplotlib.lines.Line2D at 0x12e51f61a00>,
<matplotlib.lines.Line2D at 0x12e51f6cfa0>,
<matplotlib.lines.Line2D at 0x12e51f76370>,
<matplotlib.lines.Line2D at 0x12e51f82910>,
<matplotlib.lines.Line2D at 0x12e51f82ca0>,
<matplotlib.lines.Line2D at 0x12e51f97280>,
<matplotlib.lines.Line2D at 0x12e51f97610>,
<matplotlib.lines.Line2D at 0x12e51fa4bb0>,
<matplotlib.lines.Line2D at 0x12e51fa4f40>,
<matplotlib.lines.Line2D at 0x12e51fba520>,
<matplotlib.lines.Line2D at 0x12e51fba8b0>,
<matplotlib.lines.Line2D at 0x12e51fc5e50>,
<matplotlib.lines.Line2D at 0x12e51fd0220>,
<matplotlib.lines.Line2D at 0x12e51fdc7c0>,
<matplotlib.lines.Line2D at 0x12e51fdcb50>,
<matplotlib.lines.Line2D at 0x12e51ff2130>,
<matplotlib.lines.Line2D at 0x12e51ff24c0>],
'boxes': [<matplotlib.lines.Line2D at 0x12e50de6f70>,
<matplotlib.lines.Line2D at 0x12e50e028e0>,
<matplotlib.lines.Line2D at 0x12e51f48250>,
<matplotlib.lines.Line2D at 0x12e51f56b80>,
<matplotlib.lines.Line2D at 0x12e51f6c4f0>,
<matplotlib.lines.Line2D at 0x12e51f76e20>,
<matplotlib.lines.Line2D at 0x12e51f8e790>,
<matplotlib.lines.Line2D at 0x12e51fa4100>,

```


```

<matplotlib.lines.Line2D at 0x12e51faea30>,
<matplotlib.lines.Line2D at 0x12e51fc53a0>,
<matplotlib.lines.Line2D at 0x12e51fd0cd0>,
<matplotlib.lines.Line2D at 0x12e51fe6640>],
'medians': [<matplotlib.lines.Line2D at 0x12e50e021c0>,
<matplotlib.lines.Line2D at 0x12e50e0eaf0>,
<matplotlib.lines.Line2D at 0x12e51f56460>,
<matplotlib.lines.Line2D at 0x12e51f61d90>,
<matplotlib.lines.Line2D at 0x12e51f76700>,
<matplotlib.lines.Line2D at 0x12e51f8e070>,
<matplotlib.lines.Line2D at 0x12e51f979a0>,
<matplotlib.lines.Line2D at 0x12e51fae310>,
<matplotlib.lines.Line2D at 0x12e51fbac40>,
<matplotlib.lines.Line2D at 0x12e51fd05b0>,
<matplotlib.lines.Line2D at 0x12e51fdcee0>,
<matplotlib.lines.Line2D at 0x12e51ff26a0>],
'fliers': [<matplotlib.lines.Line2D at 0x12e50e02550>,
<matplotlib.lines.Line2D at 0x12e50e0ee80>,
<matplotlib.lines.Line2D at 0x12e51f567f0>,
<matplotlib.lines.Line2D at 0x12e51f6c160>,
<matplotlib.lines.Line2D at 0x12e51f76a90>,
<matplotlib.lines.Line2D at 0x12e51f8e400>,
<matplotlib.lines.Line2D at 0x12e51f97d30>,
<matplotlib.lines.Line2D at 0x12e51fae6a0>,
<matplotlib.lines.Line2D at 0x12e51fbafd0>,
<matplotlib.lines.Line2D at 0x12e51fd0940>,
<matplotlib.lines.Line2D at 0x12e51fe62b0>,
<matplotlib.lines.Line2D at 0x12e51ff2a30>],
'means': []}

```

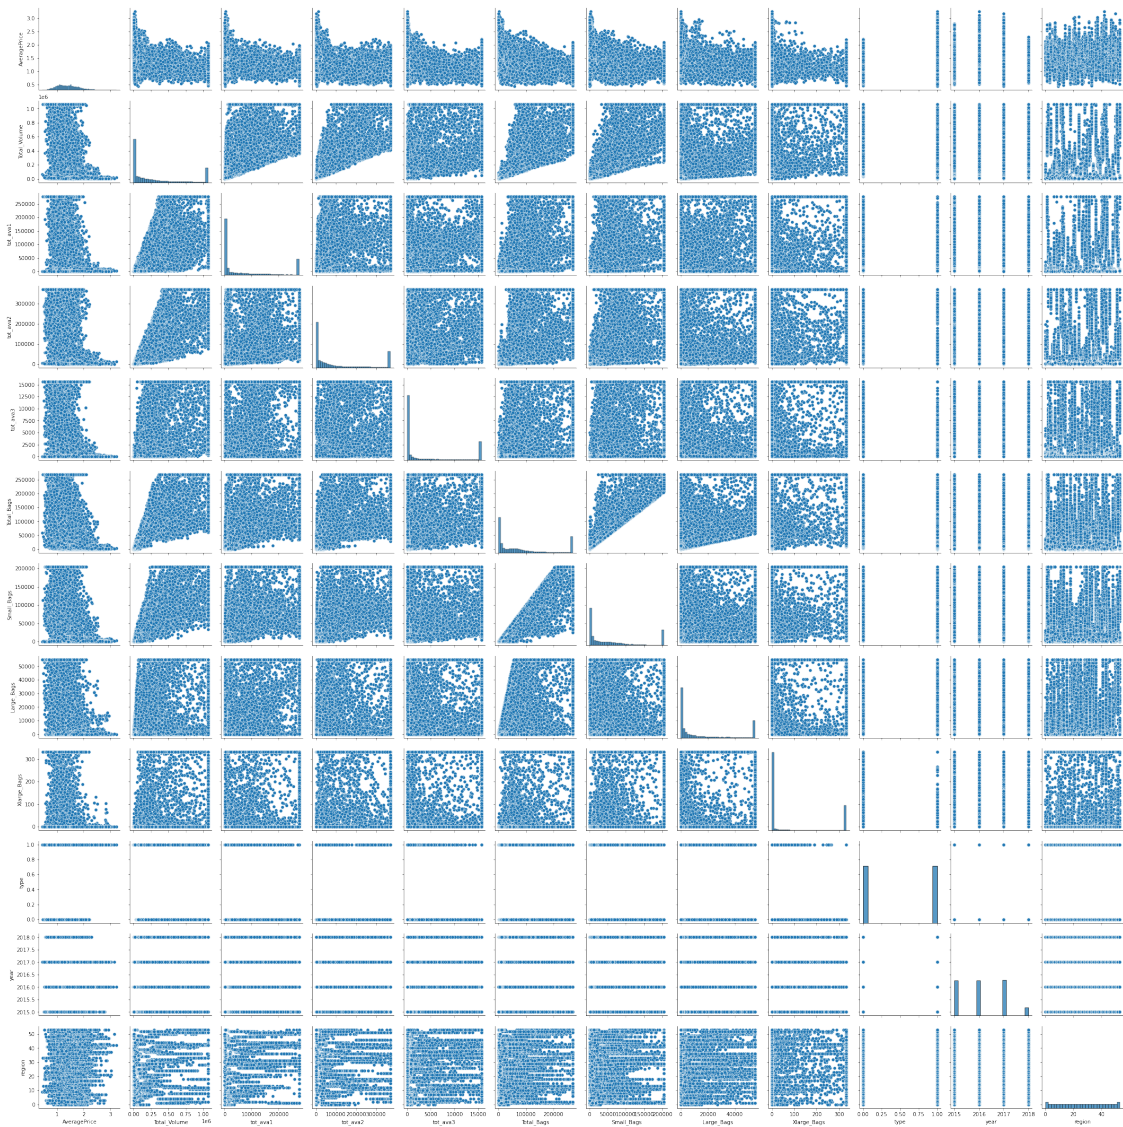



```
[142]: from scipy import stats
import pylab
stats.probplot(df.AveragePrice, dist = "norm", plot = pylab)
plt.show()
```



```
[143]: sns.pairplot(df.iloc[:, :])
```

```
[143]: <seaborn.axisgrid.PairGrid at 0x12e45f78d00>
```



```
[144]: df.corr()
```

```
[144]:
```

	AveragePrice	Total_Volume	tot_ava1	tot_ava2	tot_ava3	\
AveragePrice	1.000000	-0.503153	-0.523474	-0.412565	-0.465699	
Total_Volume	-0.503153	1.000000	0.862914	0.918696	0.735324	
tot_ava1	-0.523474	0.862914	1.000000	0.679998	0.646679	
tot_ava2	-0.412565	0.918696	0.679998	1.000000	0.716079	
tot_ava3	-0.465699	0.735324	0.646679	0.716079	1.000000	
Total_Bags	-0.486599	0.926100	0.791529	0.809947	0.663865	
Small_Bags	-0.449332	0.911307	0.759106	0.827511	0.677017	
Large_Bags	-0.424658	0.706783	0.677472	0.561624	0.435210	
Xlarge_Bags	-0.352337	0.584318	0.536757	0.553875	0.559618	
type	0.615845	-0.655466	-0.627013	-0.627230	-0.610211	

year	0.093197	0.038736	0.017481	-0.007602	-0.083827
region	-0.011716	0.106552	0.172991	0.049120	0.038485

	Total_Bags	Small_Bags	Large_Bags	Xlarge_Bags	type \
AveragePrice	-0.486599	-0.449332	-0.424658	-0.352337	0.615845
Total_Volume	0.926100	0.911307	0.706783	0.584318	-0.655466
tot_ava1	0.791529	0.759106	0.677472	0.536757	-0.627013
tot_ava2	0.809947	0.827511	0.561624	0.553875	-0.627230
tot_ava3	0.663865	0.677017	0.435210	0.559618	-0.610211
Total_Bags	1.000000	0.961362	0.775343	0.599925	-0.623950
Small_Bags	0.961362	1.000000	0.632182	0.600454	-0.620843
Large_Bags	0.775343	0.632182	1.000000	0.429142	-0.464445
Xlarge_Bags	0.599925	0.600454	0.429142	1.000000	-0.592029
type	-0.623950	-0.620843	-0.464445	-0.592029	1.000000
year	0.146330	0.128440	0.144863	0.163137	-0.000032
region	0.117088	0.112723	0.153494	0.005992	-0.000280

	year	region
AveragePrice	0.093197	-0.011716
Total_Volume	0.038736	0.106552
tot_ava1	0.017481	0.172991
tot_ava2	-0.007602	0.049120
tot_ava3	-0.083827	0.038485
Total_Bags	0.146330	0.117088
Small_Bags	0.128440	0.112723
Large_Bags	0.144863	0.153494
Xlarge_Bags	0.163137	0.005992
type	-0.000032	-0.000280
year	1.000000	-0.000055
region	-0.000055	1.000000

```
[149]: import statsmodels.formula.api as smf
```

```
[150]: model1=smf.ols("AveragePrice ~ Total_Volume + tot_ava1 + tot_ava2 + tot_ava3 +
↪Small_Bags +Large_Bags+ Xlarge_Bags+type+year+region ",data=df).fit()
```

```
[151]: model1.summary()
```

```
[151]: <class 'statsmodels.iolib.summary.Summary'>
```

```

"""
                                OLS Regression Results
=====
Dep. Variable:                AveragePrice    R-squared:                0.448
Model:                        OLS            Adj. R-squared:           0.448
Method:                        Least Squares   F-statistic:             1480.
Date:                          Mon, 23 Jan 2023  Prob (F-statistic):       0.00
Time:                          17:23:14       Log-Likelihood:          -3872.0

```

```

No. Observations:      18249   AIC:      7766.
Df Residuals:         18238   BIC:      7852.
Df Model:              10
Covariance Type:      nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      -87.8127         5.126    -17.133     0.000    -97.859    -77.766
Total_Volume  -4.499e-07     3.55e-08    -12.691     0.000   -5.19e-07   -3.8e-07
tot_ava1      -1.405e-07     5.82e-08     -2.414     0.016   -2.54e-07   -2.64e-08
tot_ava2       1.132e-06     5.6e-08     20.226     0.000     1.02e-06    1.24e-06
tot_ava3      -7.481e-06     5.94e-07    -12.596     0.000   -8.64e-06   -6.32e-06
Small_Bags     2.932e-07     8.35e-08     3.510     0.000     1.29e-07    4.57e-07
Large_Bags    -2.31e-06     1.68e-07    -13.763     0.000   -2.64e-06   -1.98e-06
Xlarge_Bags     0.0002     2.23e-05     8.440     0.000         0.000         0.000
type           0.4073         0.007     62.074     0.000         0.394         0.420
year           0.0442         0.003     17.372     0.000         0.039         0.049
region         0.0009         0.000      6.333     0.000         0.001         0.001
=====
Omnibus:                1059.748   Durbin-Watson:                0.325
Prob(Omnibus):           0.000   Jarque-Bera (JB):            1889.275
Skew:                    0.448   Prob(JB):                     0.00
Kurtosis:                4.297   Cond. No.                    1.19e+09
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.19e+09. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[152]: import statsmodels.api as sm
```

```
[153]: sm.graphics.influence_plot(model1)
```

```
[153]:
```



```
[ ]:
```

```
[87]: res_vol=smf.ols("Total_Volume ~ tot_ava1 + tot_ava2 + tot_ava3 + Total_Bags +
    ↪Small_Bags +Large_Bags+ Xlarge_Bags+type+year+region ",data=df).fit().
    ↪rsquared
```

```
[88]: res_vol = 1/(1 - res_vol)
```

```
[89]: res_vol
```

```
[89]: 39.77116230349727
```

```
[90]: res_ava1=smf.ols("tot_ava1 ~ Total_Volume + tot_ava2 + tot_ava3 + Total_Bags +
    ↪Small_Bags +Large_Bags+ Xlarge_Bags+type+year+region ",data=df).fit().
    ↪rsquared
    res_ava1 = 1/(1 - res_ava1)
```

```
[91]: res_ava1
```

```
[91]: 7.129231122219007
```

```
[92]: res_ava2=smf.ols("tot_ava2 ~ Total_Volume + tot_ava1 + tot_ava3 + Total_Bags +
    ↪Small_Bags +Large_Bags+ Xlarge_Bags+type+year+region ",data=df).fit().
    ↪rsquared
    res_ava2 = 1/(1 - res_ava2)
```

```
[93]: res_ava2
```

```
[93]: 12.281073160797487
```

```
[94]: res_ava3=smf.ols("tot_ava3 ~ Total_Volume + tot_ava1 + tot_ava2 + Total_Bags +
    ↪Small_Bags +Large_Bags+ Xlarge_Bags+type+year+region ",data=df).fit().
    ↪rsquared
    res_ava3 = 1/(1 - res_ava3)
```

```
[95]: res_ava3
```

```
[95]: 2.5860951329231647
```

```
[96]: res_Total_Bags=smf.ols("Total_Bags ~ Total_Volume + tot_ava1 + tot_ava2 +
    ↪tot_ava3 + Small_Bags +Large_Bags+ Xlarge_Bags+type+year+region ",data=df).
    ↪fit().rsquared
    res_Total_Bags = 1/(1 - res_Total_Bags)
```

```
[97]: res_Total_Bags
```

[97]: 39.88789871182136

```
[101]: res_Small_Bags=smf.ols("Total_Bags ~ Total_Volume + tot_ava1 + tot_ava2 +  
      ↪tot_ava3 + Total_Bags +Large_Bags+ Xlarge_Bags+type+year+region ",data=df).  
      ↪fit().rsquared  
      res_Small_Bags = 1/(1 - res_Small_Bags)
```

C:\Users\Ravi\AppData\Local\Temp\ipykernel_12936\662951263.py:2: RuntimeWarning:
divide by zero encountered in double_scalars
 res_Small_Bags = 1/(1 - res_Small_Bags)

```
[102]: res_Small_Bags
```

[102]: inf

```
[103]: res_Large_Bags=smf.ols("Large_Bags ~ Total_Volume + tot_ava1 + tot_ava2 +  
      ↪tot_ava3 + Total_Bags +Small_Bags+ Xlarge_Bags+type+year+region ",data=df).  
      ↪fit().rsquared  
      res_Large_Bags = 1/(1 - res_Large_Bags)
```

```
[104]: res_Large_Bags
```

[104]: 4.693747595611235

```
[105]: res_Xlarge_Bags=smf.ols("Xlarge_Bags ~ Total_Volume + tot_ava1 + tot_ava2 +  
      ↪tot_ava3 + Total_Bags +Small_Bags+ Large_Bags+type+year+region ",data=df).  
      ↪fit().rsquared  
      res_Xlarge_Bags = 1/(1 - res_Xlarge_Bags)
```

```
[106]: res_Xlarge_Bags
```

[106]: 1.9457175245752574

```
[107]: res_type=smf.ols("type ~ Total_Volume + tot_ava1 + tot_ava2 + tot_ava3 +  
      ↪Total_Bags +Small_Bags+ Large_Bags+Xlarge_Bags+year+region ",data=df).fit().  
      ↪rsquared  
      res_type = 1/(1 - res_type)
```

```
[108]: res_type
```

[108]: 2.197432470456866

```
[109]: res_year=smf.ols("year ~ Total_Volume + tot_ava1 + tot_ava2 + tot_ava3 +  
      ↪Total_Bags +Small_Bags+ Large_Bags+Xlarge_Bags+type+region ",data=df).fit().  
      ↪rsquared  
      res_year = 1/(1 - res_year)
```

```
[110]: res_year
```



```
[110]: 1.1677948547468082
```

```
[111]: res_region=smf.ols("region ~ Total_Volume + tot_ava1 + tot_ava2 + tot_ava3 +  
    ↪Total_Bags +Small_Bags+ Large_Bags+Xlarge_Bags+type+year ",data=df).fit().  
    ↪rsquared  
res_region = 1/(1 - res_region)
```

```
[112]: res_region
```

```
[112]: 1.0767857251331638
```

```
[113]: df1 = {'Variables':['region','Total_Volume', 'tot_ava1',  
    ↪'tot_ava2','tot_ava3','Total_Bags','Small_Bags','Large_Bags','Xlarge_Bags','type','year'],  
    ↪'VIF':[res_region,res_vol, res_ava1,  
    ↪res_ava2,res_ava3,res_Total_Bags,res_Small_Bags,res_Large_Bags,res_Xlarge_Bags,res_type,res  
Vif_frame = pd.DataFrame(df1)
```

```
[114]: Vif_frame
```

```
[114]:
```

	Variables	VIF
0	region	1.076786
1	Total_Volume	39.771162
2	tot_ava1	7.129231
3	tot_ava2	12.281073
4	tot_ava3	2.586095
5	Total_Bags	39.887899
6	Small_Bags	inf
7	Large_Bags	4.693748
8	Xlarge_Bags	1.945718
9	type	2.197432
10	year	1.167795

```
[158]: model2=smf.ols("AveragePrice ~ tot_ava1 + tot_ava3 + Small_Bags +Large_Bags+  
    ↪Xlarge_Bags+type+year+region ",data=df).fit()
```

```
[159]: model2.summary()
```

```
[159]: <class 'statsmodels.iolib.summary.Summary'>  
      ""
```

```

                        OLS Regression Results
=====
Dep. Variable:          AveragePrice   R-squared:                0.434
Model:                  OLS           Adj. R-squared:            0.434
Method:                 Least Squares   F-statistic:              1751.
Date:                  Mon, 23 Jan 2023   Prob (F-statistic):       0.00
Time:                  17:39:58          Log-Likelihood:           -4093.6
No. Observations:      18249            AIC:                     8205.
Df Residuals:          18240            BIC:                     8276.
```

```

Df Model:                8
Covariance Type:         nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      -76.7944         5.137     -14.948      0.000     -86.864     -66.725
tot_ava1      -6.904e-07      4.1e-08     -16.835      0.000     -7.71e-07     -6.1e-07
tot_ava3      -5.266e-06      5.74e-07     -9.178      0.000     -6.39e-06     -4.14e-06
Small_Bags      3.227e-07      5.81e-08      5.553      0.000      2.09e-07      4.37e-07
Large_Bags     -2.808e-06      1.58e-07    -17.768      0.000     -3.12e-06     -2.5e-06
Xlarge_Bags      0.0002      2.25e-05      9.950      0.000      0.000      0.000
type           0.3832         0.007     58.600      0.000      0.370      0.396
year           0.0387         0.003     15.192      0.000      0.034      0.044
region         0.0009         0.000      6.232      0.000      0.001      0.001
=====
Omnibus:                986.868   Durbin-Watson:                0.314
Prob(Omnibus):           0.000   Jarque-Bera (JB):        1685.127
Skew:                    0.436   Prob(JB):                 0.00
Kurtosis:                4.207   Cond. No.                3.37e+08
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.37e+08. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[160]: pre1=model2.predict(df)
```

```
[161]: res1 = df.AveragePrice - pre1
res_sqr1 = res1 * res1
mse1 = np.mean(res_sqr1)
rmse1 = np.sqrt(mse1)
rmse1
```

```
[161]: 0.30281947789261016
```

```
[162]: model3=smf.ols("AveragePrice ~ np.log(tot_ava1 + tot_ava3 + Small_Bags_
↪+Large_Bags+ Xlarge_Bags+type+year+region)",data=df).fit()
```

```
[163]: model3.summary()
```

```
[163]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

=====

```

Dep. Variable:          AveragePrice    R-squared:                0.379
Model:                  OLS             Adj. R-squared:           0.379
Method:                 Least Squares   F-statistic:              1.113e+04
Date:                   Mon, 23 Jan 2023 Prob (F-statistic):       0.00
Time:                   17:40:44        Log-Likelihood:           -4948.3
No. Observations:      18249           AIC:                     9901.
Df Residuals:          18247           BIC:                     9916.
Df Model:               1
Covariance Type:       nonrobust

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept
2.9384      0.015     199.710      0.000      2.910      2.967
np.log(tot_ava1 + tot_ava3 + Small_Bags + Large_Bags + Xlarge_Bags + type + year
+ region)  -0.1422      0.001    -105.507      0.000     -0.145     -0.140
=====
Omnibus:                 1311.280    Durbin-Watson:                0.252
Prob(Omnibus):            0.000    Jarque-Bera (JB):            1798.659
Skew:                     0.624    Prob(JB):                     0.00
Kurtosis:                 3.900    Cond. No.                     69.0
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

```

[164]: pre2=model3.predict(df)

res2 = df.AveragePrice - pre2
res_sqr2 = res2 * res2
mse2 = np.mean(res_sqr2)
rmse2 = np.sqrt(mse2)
rmse2

```

```

[164]: 0.31733969488554886

```

```

[165]: model4=smf.ols("np.log(AveragePrice) ~ (tot_ava1 + tot_ava3 + Small_Bags_
↪+Large_Bags+ Xlarge_Bags+type+year+region) ",data=df).fit()

```

```

[166]: model4.summary()

```

```

[166]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

OLS Regression Results

```

=====
Dep. Variable:      np.log(AveragePrice)    R-squared:                0.457
Model:              OLS                    Adj. R-squared:           0.456
Method:             Least Squares          F-statistic:             1916.
Date:              Mon, 23 Jan 2023        Prob (F-statistic):       0.00
Time:              17:49:03                Log-Likelihood:          2282.1
No. Observations:   18249                  AIC:                     -4546.
Df Residuals:       18240                  BIC:                     -4476.
Df Model:           8
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-50.0463	3.622	-13.816	0.000	-57.147	-42.946
tot_ava1	-6.849e-07	2.89e-08	-23.685	0.000	-7.42e-07	-6.28e-07
tot_ava3	-5.098e-06	4.05e-07	-12.601	0.000	-5.89e-06	-4.31e-06
Small_Bags	3.84e-07	4.1e-08	9.370	0.000	3.04e-07	4.64e-07
Large_Bags	-2.12e-06	1.11e-07	-19.030	0.000	-2.34e-06	-1.9e-06
Xlarge_Bags	0.0002	1.58e-05	11.829	0.000	0.000	0.000
type	0.2590	0.005	56.170	0.000	0.250	0.268
year	0.0249	0.002	13.873	0.000	0.021	0.028
region	0.0007	0.000	6.542	0.000	0.000	0.001

```

=====
Omnibus:            613.278    Durbin-Watson:           0.344
Prob(Omnibus):      0.000     Jarque-Bera (JB):         925.240
Skew:               -0.331    Prob(JB):                 1.22e-201
Kurtosis:           3.882     Cond. No.                  3.37e+08
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.37e+08. This might indicate that there are strong multicollinearity or other numerical problems.

"""

[]:

```

[167]: pre3=model4.predict(df)

res3 = df.AveragePrice - pre3
res_sqr3 = res3 * res3
mse3 = np.mean(res_sqr3)
rmse3 = np.sqrt(mse3)
rmse3

```

```
[167]: 1.1493230854794034
```

```
[137]: data = {"MODEL":pd.Series(["model2", "Log model", "Exp model"]), "RMSE":pd.
        ↪Series([rmse1, rmse2, rmse3])}
        table_rmse = pd.DataFrame(data)
        table_rmse
```

```
[137]:
```

	MODEL	RMSE
0	model2	0.302819
1	Log model	0.317340
2	Exp model	1.149323

```
[168]: final_model=smf.ols("(AveragePrice) ~ (tot_ava1 + tot_ava3 + Small_Bags_
        ↪+Large_Bags+ Xlarge_Bags+type+year+region) ",data=df).fit()
```

```
[169]: final_model.summary()
```

```
[169]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:                AveragePrice    R-squared:                0.434
Model:                        OLS            Adj. R-squared:         0.434
Method:                      Least Squares    F-statistic:              1751.
Date:                        Mon, 23 Jan 2023    Prob (F-statistic):       0.00
Time:                        17:52:20          Log-Likelihood:           -4093.6
No. Observations:            18249            AIC:                     8205.
Df Residuals:                18240            BIC:                     8276.
Df Model:                    8
Covariance Type:              nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
Intercept                -76.7944         5.137    -14.948     0.000    -86.864    -66.725
tot_ava1                 -6.904e-07     4.1e-08    -16.835     0.000    -7.71e-07    -6.1e-07
tot_ava3                 -5.266e-06     5.74e-07    -9.178     0.000    -6.39e-06    -4.14e-06
Small_Bags               3.227e-07     5.81e-08     5.553     0.000     2.09e-07     4.37e-07
Large_Bags              -2.808e-06     1.58e-07   -17.768     0.000    -3.12e-06    -2.5e-06
Xlarge_Bags               0.0002     2.25e-05     9.950     0.000         0.000         0.000
type                     0.3832         0.007     58.600     0.000         0.370         0.396
year                     0.0387         0.003     15.192     0.000         0.034         0.044
region                   0.0009         0.000      6.232     0.000         0.001         0.001
=====
Omnibus:                    986.868    Durbin-Watson:           0.314
Prob(Omnibus):              0.000    Jarque-Bera (JB):        1685.127
Skew:                      0.436    Prob(JB):                 0.00
Kurtosis:                   4.207    Cond. No.                 3.37e+08
=====
```

=====
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.37e+08. This might indicate that there are strong multicollinearity or other numerical problems.

"""


```
pre_final=final_model.predict(df)
```

```
[170]: pre_final
```

```
-----  
NameError                                Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_18108\2097908561.py in <module>  
----> 1 pre_final  
  
NameError: name 'pre_final' is not defined
```

```
[171]: res = final_model.resid  
sm.qqplot(res)  
plt.show()
```


```
C:\Users\Ravi\anaconda3\lib\site-packages\statsmodels\graphics\gofplots.py:993:  
UserWarning: marker is redundantly defined by the 'marker' keyword argument and  
the fmt string "bo" (-> marker='o'). The keyword argument will take precedence.  
ax.plot(x, y, fmt, **plot_style)
```




```
[144]: pre_final
```

```
[144]: 0      1.218806
      1      1.219250
      2      1.218331
      3      1.217554
      4      1.217588
      ...
     18244    1.767211
     18245    1.766770
     18246    1.763862
     18247    1.764120
     18248    1.766234
     Length: 18249, dtype: float64
```

```
[172]: # Q-Q plot
stats.probplot(res, dist = "norm", plot = pylab)
plt.show()
```





```
[146]: sns.residplot(x = pred, y = df.AveragePrice, lowess = True)
plt.xlabel('Fitted')
plt.ylabel('Residual')
plt.title('Fitted vs Residual')
plt.show()
```

```
[122]: sm.graphics.influence_plot(final_model)
```

```
[122]:
```





```
[173]: from sklearn.model_selection import train_test_split
df_train, df_test = train_test_split(df, test_size = 0.2, random_state=0) # 20%
      ↪ test data
```

```
[174]: model_train = smf.ols("AveragePrice ~ tot_ava1 + tot_ava3 + Total_Bags +
      ↪ Small_Bags + Large_Bags + Xlarge_Bags + type + year + region", data = df_train).fit()
```

```
[175]: test_pred = model_train.predict(df_test)
```

```
[176]: # test residual values
test_resid = test_pred - df_test.AveragePrice
```

```
[177]: # RMSE value for test data
test_rmse = np.sqrt(np.mean(test_resid * test_resid))
test_rmse
```

```
[177]: 0.3025386373302603
```

```
[178]: # train_data prediction
train_pred = model_train.predict(df_train)
```

```
[179]: # train residual values
train_resid = train_pred - df_train.AveragePrice
```

```
[180]: # RMSE value for train data
train_rmse = np.sqrt(np.mean(train_resid * train_resid))
train_rmse
```

```
[180]: 0.3015786165917111
```

```
[181]: train_rmse
```

```
[181]: 0.3015786165917111
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```